

BACHELOR INFORMATICA



UNIVERSITY OF AMSTERDAM

CLIsis: An Interface for Visually Impaired Users of Apache Isis Applications

Sander Ginn

May 21, 2016

Supervisor(s): Maarten Marx (UvA), Dan Haywood (Apache)

Signed: Maarten Marx (UvA)

Abstract

This will be the abstract.

Contents

1	Introduction	4
1.1	Research questions	6
1.2	Overview of thesis	7
2	Theoretical background	8
2.1	Apache Isis	9
2.1.1	Domain-Driven Design	10
2.1.2	Naked Objects pattern	11
2.2	User interface migration in history	12
3	Methods	13
3.1	User interface migration	15
3.1.1	Detection	15
3.1.2	Representation	15
3.1.3	Transformation	15
3.2	Experiments	15
3.2.1	GOMS	15
3.2.2	Time trial	15
4	Implementation	16
4.1	Frameworks and tools	17
4.1.1	REST API	17
4.1.2	AngularJS	17
4.1.3	Web Speech API	17
4.2	Functionality	17
4.2.1	Main module	17
4.2.2	Controllers	17
4.2.3	Views	17
5	Evaluation	18
5.1	Specification requirements	19
5.2	Experiments	19
5.2.1	GOMS	19
5.2.2	Time trial	19
5.2.3	Performance difference	19
6	Conclusion	20
6.1	Future work	21
6.2	Acknowledgements	21
A	Application manual	28
A.1	Installation	29
A.2	User manual	29
B	Time trial test class	30

Introduction

Every application that offers a method of interaction with a human being requires some form of user interface to enable this interaction. But what exactly characterises a user interface? The Oxford English Dictionary provides the following definition for "interface" [1]:

in·ter·face, *n.*: A means or place of interaction between two systems, organizations, etc.; a meeting-point or common ground between two parties, systems, or disciplines; also, interaction, liaison, dialogue.

Thus, in the context of an application, a user interface is a means of interaction between the user and system based on common ground between these two entities. The common ground is a layer on top of the system which presents information from the system in a comprehensible manner to the user and translates user input to system-interpretable commands.

This is still a very abstract definition of what a user interface does. It does not specify how system information is presented, how the user provides input, what permissions and restrictions the interface must adhere to, et cetera. Nielsen[2] describes an overview of user interfaces throughout history which will provide some practical examples.

Batch systems were one of the first user interfaces to be introduced. To interact with the system, a user submitted all of the system commands in one batch and had to wait for all of them to be processed before being presented with the results. It was a tedious and inconvenient way of interacting with systems, as erroneous commands would prevent the batch from successfully completing. Early batch systems also often required special operators to program the commands. Contemporary user interfaces generally still offer batch operation in some form or another, as its main advantage is that it can be used to perform menial tasks without continuous human supervision.

As the use of computers became more widespread, the *line-oriented interface* was adopted as a method of interaction. The name is derived from the characteristic that the user interacted with a system on a single line; once submitted, the command could not be modified anymore. This characteristic enforces a use pattern where a dialog between the system and user takes place. The user answers questions posed by the system and issues commands with parameters. As long as the information is well structured this method of interaction still has valid use cases in present times, as it effectively delimits functionality, preventing novice users from running into issues.

The successor of line-oriented interfaces is the *full-screen interface*. This interface aims to exploit a greater deal of the screen real estate by enabling the user to interact with more elements than just one input line, such as in the interface in figure 1.1. The full-screen interface also introduced the concept of menus, where functionality available to the user is listed in a hierarchical structure. Menu design has been researched extensively[3, 4, 5] to determine the optimal balance of depth and breadth. Depth decreases complexity of menus while increasing ease of navigation; the contrary applies to breadth. Full-screen interfaces are still used present-day (e.g. electronic forms), albeit in adapted forms such as pop-up windows.

Finally, the user interface type which is currently the most widespread, is the *graphical user interface* (GUI). The vast majority of modern applications offer a graphical user interface as their primary method of interaction. A key characteristic of GUIs is that interaction is offered through direct manipulation. Rather than issuing a command which exactly specifies those parameters of what has to happen, the user gets continuous feedback of their input commands. An example is resizing a window: if a user were to do this by a command, the dimensions are

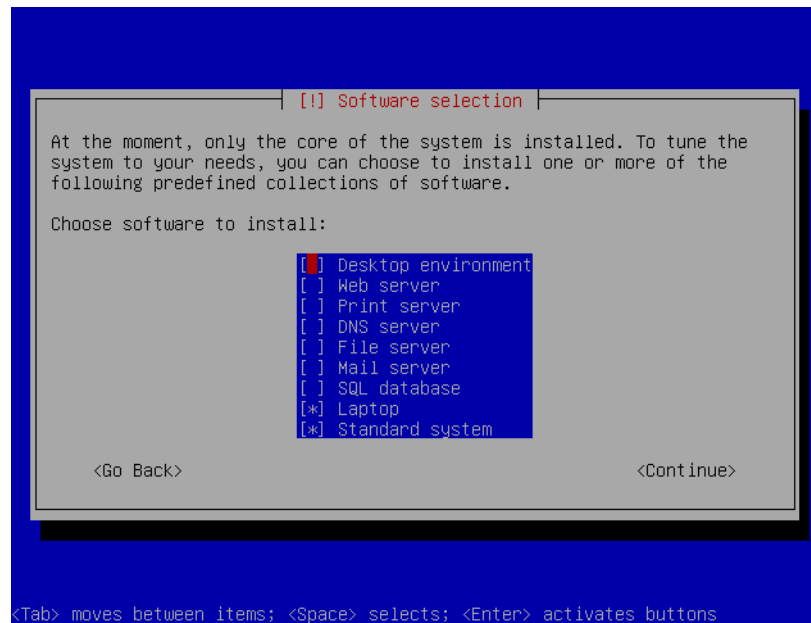


Figure 1.1: A screenshot of Debian’s full-screen interface during installation

provided as parameters and the window resizes. In a GUI, the user can dynamically resize the window with the mouse. The obvious advantage of a GUI is that it can take advantage of the human sight to represent data in a more intuitive way. Research has shown that when the interface adheres to some fundamental principles, a user’s cognitive performance increases with a GUI. An example of one of these principles is Miller’s Law, proposed in his seminal research on information processing[6]. However, Nielsen also addresses the issue that motivates this research: a GUI sidelines the visually impaired from using an application.

Enabling the visually impaired to use applications despite a lack of or very bad vision has been an ongoing effort since GUIs were widely accepted as the de facto standard[7]. However, it has proven to be difficult to design a GUI that is also accessible without visual cues. A big issue is that standardisation has proven ineffective: even though the United States has passed laws that enforce accessibility of websites[8] and the World Wide Web Consortium has developed the Web Content Accessibility Guidelines[9], research has shown that these guidelines are very often ignored and thus websites are rendered inaccessible for accessibility tools such as screenreaders[10].

This research aims to provide an alternative interface for visually impaired users for applications developed with the Apache Isis[11] framework. This interface will not serve as a replacement for the existing interface nor blend in with it. The interface is derived from the metamodel of the framework and thus content independent. This means that the interface is readily available for any Isis application and no alterations are necessary to add it to an existing application.

1.1 Research questions

The central research question in this thesis is:

Can a content independent graphical user interface be adapted to a non-visual interface so that visually impaired users are able to interact with the application?

For the non-visual interface to be useful it is imperative that it offers the same functionality as the graphical user interface. Therefore, we will also answer the following subquestion:

When implementing a non-visual interface, can the integrity of the domain model be maintained while providing an effective and simple method of interaction?

Finally, it is desired that the non-visual interface does not severely impact user performance, as this would imply that it is not a legitimate alternative to the existing interface. Thus, a second subquestion will be answered:

Compared to the standard interface, how is performance affected when the user employs the new interface?

1.2 Overview of thesis

A short overview of what this thesis will describe is given.

Theoretical background

User interface design is a thoroughly studied discipline with strong roots in psychology. In the 1980s GUI development exploded due to better hardware[12]. This meant that traditional user interfaces had to be redesigned to accommodate to the graphical features of the modern computer. In section 2.2 we will provide a brief history on how this was achieved and what sort of issues arose when migrating a user interface. Furthermore, section 2.1 will describe what Apache Isis entails.

2.1 Apache Isis

As a software engineer, picking a framework for developing web applications can be a tedious process. There are dozens of frameworks for Java alone, with the oldest and most adopted one being Spring[13]. The vast majority of these frameworks are based on the *model-view-controller* (MVC) pattern, where the view displays information to the user, the controller processes interaction between the user and the view, and the model contains the information and logic that manipulates this information[14]. The relations between the components are depicted in figure 2.1.

While the MVC pattern itself has a lot of advantages, it has received criticism in the context of web development. The concept of separating business logic from presentation logic is often not adhered to in web applications, resulting in controllers that are contaminated with logic that should live in the model[15].




Figure 2.1: Interaction between the components of an MVC application

This is where Apache Isis differs from more classic MVC-based frameworks. There is no need for an explicit controller layer, as the core strength of the framework is that behaviour which would normally be defined in controllers is automatically generated by reflection¹. A major benefit of this feature is that prototyping is greatly simplified; the business logic is virtually all that needs to be programmed, and the framework does the rest. This means that project managers, product owners and other parties involved can be presented with working prototypes rather than charts and specifications, and it relieves developers from wasting precious time on coding the visible part of an application - code which often has limited reusability in a later stage of development.

This does not, however, imply that Apache Isis is merely useful for prototyping. The automatically generated interface is perfectly suitable for a lot of business oriented applications.

¹Reflection is a language's ability to inspect and dynamically call classes, methods, attributes, etc. at runtime.



Figure 2.2: A very simple example of how the user interface is derived from code

Estatio[16], an estate management application built on Apache Isis commissioned by a listed real estate investment firm, makes use of the standard wicket viewer. In case a given application demands a bespoke user interface, it can be designed to use the REST API which ships with Apache Isis, which we will use to implement our interface. Furthermore, there are nearly thirty add-ons available which can be plugged into any Apache Isis application[17], which offer various features such as Excel spreadsheet importing. Finally, the framework takes testing support very seriously, offering unit and integration testing functionality to ensure the business logic behaves as expected.

Apache Isis is developed on two fundamental philosophies: domain-driven design and the Naked Objects pattern.

2.1.1 Domain-Driven Design

Software projects have been riddled by failure for as long as they have been around. A 1999 survey by Linberg found that 20% of software projects failed and another 46% suffered from severe delays, exceeding budgets or missing functionality or a combination of the three[18]. In 2007, the Standish Group concluded that 35% of software projects could be classified as successful, with 46% suffering from the aforementioned problems and 19% completely failing[19]. Keil et al. concluded that a key factors in the failure of software projects are information asymmetry and goal incongruence[20].

Those exact issues are at the core of what *domain-driven design* (DDD), conceived by Eric Evans in 2004, encapsulates. In a business software project, there are two main parties involved: the domain experts (the business experts) and the technical experts (the developers). These parties will have to work together in order to specify the functionality of the software. Business language, however, is very different from technical language, and thus some method of translation is needed. A business analyst could provide this translation, but this only adds another layer of potential misunderstanding, ultimately culminating in faulty behaviour being implemented.

Ubiquitous language

The first central concept in DDD is to create an *ubiquitous language* between the domain experts and technical experts. By clearly defining what terms to use for what concepts and ensuring that both parties understand what they mean, information asymmetry can be avoided[21]. These defined terms together form the domain model which is used to describe and specify the software. Undoubtedly, at some stage in the process the ubiquitous language will prove itself insufficient to accurately describe new functionality, which implies that the model should be expanded. This creates an iterative development process where both parties have to put in effort to be as concise as possible. It also prevents the domain model from growing too complex: if the parties fail to define a certain aspect of the model, it is highly unlikely that it will be successfully implemented in the software and it is better to leave it out.

Model-driven design

Once the domain model has been specified to a sufficient extent to initiate the development process, the second central concept in DDD comes to light: *model-driven design*. As the name suggests, this means writing your software driven by the domain model. The domain model will function as the translation layer between the domain experts and the technical experts, as the software becomes more comprehensible for the domain experts if the software uses the same ubiquitous language. The code and domain model share a bidirectional relationship: if the domain model is extended, the code must be changed, and if the code is changed, the domain model must be adapted[21]. This prevents the issue of goal incongruence, as the domain experts and technical experts are aligned in terms of what has to be implemented. After all, the domain model is binding and defined in a collaborative effort.

2.1.2 Naked Objects pattern

One of the frustrations often expressed regarding DDD is that while the ubiquitous language combined with the derived domain model may very well tackle the problem of goal diversion, it also increases overhead greatly. Consider figure 2.3: the domain model is represented in the business logic layer. Recall that we've stated that the domain model will be refined over the course of time, and that any changes to the domain model must be reflected in the code. This means that any modifications to the domain model will have to be applied to the other layers, too.




Figure 2.3: The main application layers in DDD

To address this issue, Richard Pawson designed the Naked Objects pattern, based on three core concepts[22]:

1. Business logic should only reside in domain objects²
2. The user interface is an exact representation of the domain objects, and thus is an object-oriented user interface
3. The user interface is derived completely and fully automatically from the domain objects through reflection

These three concepts combined provide a means of developing software that robustly follows the philosophy of DDD. The first two concepts ensure that the software implements the domain model and prevents the use of technical language. The application, whether prototypal or deployed, will feel familiar to the stakeholders and developers alike. The third concept is not necessarily related to DDD, but does take away the burden of developing the presentation layer and thus places the focus on the business logic: i.e., the domain model.

²A domain object is an entity that holds relevant data, in order to transfer the data between the application layers.

2.2 User interface migration in history

In the history of user interfaces, two events stand out: the migration from character-based interfaces to graphical interfaces, and implementing a web-based interface for a native application.

As the use of GUIs became widespread, a lot of older software still made use of legacy user interfaces. These interfaces were usually character-based interfaces, developed for hardware that did not offer the capabilities to use a GUI[23]. Developers quickly came to notice that the migration of one user interface to the other was not easily accomplished, as end users wanted novel user interfaces to have identical functionality as the legacy counterpart, and a lack of homogenisation in hardware support resulted in a vast amount of code refactoring to achieve the desired result[24]. On top of that, simply redesigning the software from scratch would be an effort too costly in terms of resources and investments. As a result, the developers often reverse engineered the system, creating an object-oriented model and then wrapping the legacy objects so that the new system could use these entities[25]. This method allowed developers to incrementally phase out parts of the software until all legacy objects were removed.

Another major shift in user interface design came to the forefront when the internet was adopted in offices and homes alike and applications started to become available in the web browser rather than a native application for independent operating systems. This migration is still taking place to this day; an example is Skype, of which Microsoft introduced a preview version of its web based version as late 2013[26]. Developers now had to deal with writing the presentation layer in all-new web technologies which run client side, while respecting the many constraints dictated by the software. When the underlying logic is not too complex, developers usually opted for "grafting" the web interface on top of the application, designing most of the mutually independent functionalities from scratch and applying reverse engineering where necessary. For more complex applications, however, a comparable wrapping method as mentioned before could also be employed[27].

Developers often also face the common issue of "dead" or "glue" code when migrating legacy systems to a web version, requiring the developers to thoroughly analyse code to prevent obsolete functionality from being implemented in the new user interface. Moreover, partially reusing code modules may also violate application logic[28]. Finally, the great advantage of web availability of software is that it becomes accessible to a vast group of users: e.g. a bank offering web-based banking to its customers, removing the need of trained bank clerks. This does, however, introduce another layer of complexity, as the various groups of users should interact with the user interfaces designed for them - corporate clients require more functionality than a personal customer, but they both speak to the same back end[27].

Methods

This chapter will describe the employed methods.

3.1 User interface migration

This section will describe what method I have applied to perform the user interface migration.

3.1.1 Detection

Step one of the migration method.

3.1.2 Representation

Step two of the migration method.

Specifications

The set of specifications that are derived from section 3.1.1 are presented.

3.1.3 Transformation

Step three of the migration method.

3.2 Experiments

A description of the methods used to conduct the experiments.

3.2.1 GOMS

The Goals, Operators, Methods, Selection rules method is explained.

3.2.2 Time trial

A custom time trial method is described.

Implementation

This chapter describes the practical side of the project.

4.1 Frameworks and tools

This section expands on what available software tools I have used to implement the new interface.

4.1.1 REST API

This subsection will describe how Apache Isis' REST API works and how it is incorporated in the new interface.

4.1.2 AngularJS

Here I will explain the choice for AngularJS.

4.1.3 Web Speech API

This subsection will describe how the Web Speech API is used.

4.2 Functionality

This section will expand on the individual components that process the information from the REST API.

4.2.1 Main module

This subsection describes the main app module.

4.2.2 Controllers

This subsection describes each controller's functionality.

4.2.3 Views

This subsection describes the purpose of each view.

Evaluation

This chapter describes the experiment results and answers the research questions.

5.1 Specification requirements

In this section, I will motivate which specifications have been met in the implementation. With this information I will be able to answer the first subquestion.

5.2 Experiments

5.2.1 GOMS

Here, the results from applying GOMS will be evaluated.

5.2.2 Time trial

Here, the results from the time trial experiment will be evaluated.

5.2.3 Performance difference

This subsection will evaluate the experiment results and answer the second subquestion.

Conclusion

Here, I will answer the research question and make any other conclusions.

6.1 Future work

Any remaining future work can be discussed here.

6.2 Acknowledgements

Bibliography

- [1] “interface, n. : Oxford english dictionary.” <http://www.oed.com/view/Entry/97747?result=1&rskey=X7p96Q&>. (Accessed on 05/16/2016).
- [2] J. Nielsen, *Usability engineering*. Elsevier, 1994.
- [3] K. R. Paap and R. J. Roske-Hofstrand, “The optimal number of menu options per panel,” *Human Factors: The Journal of the Human Factors and Ergonomics Society*, vol. 28, no. 4, pp. 377–385, 1986.
- [4] T. K. Landauer and D. Nachbar, “Selection from alphabetic and numeric menu trees using a touch screen: breadth, depth, and width,” *ACM SIGCHI Bulletin*, vol. 16, no. 4, pp. 73–78, 1985.
- [5] D. L. Fisher, E. J. Yungkurth, and S. M. Moss, “Optimal menu hierarchy design: syntax and semantics,” *Human Factors: The Journal of the Human Factors and Ergonomics Society*, vol. 32, no. 6, pp. 665–683, 1990.
- [6] G. A. Miller, “The magical number seven, plus or minus two: some limits on our capacity for processing information,” *Psychological review*, vol. 63, no. 2, p. 81, 1956.
- [7] L. H. Boyd *et al.*, “The graphical user interface crisis: Danger and opportunity,” 1990.
- [8] “Section 508 of the rehabilitation act.” <http://www.section508.gov/section-508-of-the-rehabilitation-act>. (Accessed on 05/16/2016).
- [9] “Web content accessibility guidelines (wcag) 2.0.” <https://www.w3.org/TR/WCAG20/>. (Accessed on 05/16/2016).
- [10] S. Leuthold, J. A. Bargas-Avila, and K. Opwis, “Beyond web content accessibility guidelines: Design of enhanced text user interfaces for blind internet users,” *International Journal of Human-Computer Studies*, vol. 66, no. 4, pp. 257–270, 2008.
- [11] “Apache isis.” <http://isis.apache.org/>. (Accessed on 05/16/2016).
- [12] B. A. Myers, “A brief history of human-computer interaction technology,” *interactions*, vol. 5, no. 2, pp. 44–54, 1998.
- [13] “Spring.” <https://spring.io/>. (Accessed on 05/16/2016).
- [14] A. Leff and J. T. Rayfield, “Web-application development using the model/view/controller design pattern,” in *Enterprise Distributed Object Computing Conference, 2001. EDOC’01. Proceedings. Fifth IEEE International*, pp. 118–127, IEEE, 2001.
- [15] “Fulfilling the promise of mvc.” <https://www.infoq.com/articles/Nacked-MVC>. (Accessed on 05/17/2016).
- [16] “Estatio.” <http://www.estatio.org/>. (Accessed on 05/18/2016).
- [17] “Apache isis add-ons.” <http://www.isisaddons.org/>. (Accessed on 05/18/2016).
- [18] K. R. Linberg, “Software developer perceptions about software project failure: a case study,” *Journal of Systems and Software*, vol. 49, no. 2, pp. 177–192, 1999.
- [19] D. Rubinstein, “Standish group report: Theres less development chaos today,” *Software Development Times*, vol. 1, 2007.

- [20] M. Keil, J. Mann, and A. Rai, “Why software projects escalate: An empirical analysis and test of four theoretical models,” *Mis Quarterly*, pp. 631–664, 2000.
- [21] E. Evans, *Domain-driven design: tackling complexity in the heart of software*. Addison-Wesley Professional, 2004.
- [22] R. Pawson and R. Matthews, “Naked objects,” in *Companion of the 17th annual ACM SIG-PLAN conference on Object-oriented programming, systems, languages, and applications*, pp. 36–37, ACM, 2002.
- [23] E. Merlo, P.-Y. Gagné, J.-F. Girard, K. Kontogiannis, L. Hendren, P. Panangaden, and R. De Mori, “Reengineering user interfaces,” *IEEE Software*, vol. 12, no. 1, p. 64, 1995.
- [24] M. Moore, S. Rugaber, and P. Seaver, “Knowledge-based user interface migration,” in *Software Maintenance, 1994. Proceedings., International Conference on*, pp. 72–79, IEEE, 1994.
- [25] A. De Lucia, G. Lucca, A. R. Fasolino, P. Guerra, and S. Petruzzelli, “Migrating legacy systems towards object-oriented platforms,” in *Software Maintenance, 1997. Proceedings., International Conference on*, pp. 122–129, IEEE, 1997.
- [26] “Type less. talk more. make skype calls directly from your inbox.” <http://blogs.skype.com/2013/04/29/type-less-talk-more-make-skype-calls-directly-from-your-outlook-com-inbox>. (Accessed on 05/21/2016).
- [27] L. Aversano, G. Canfora, A. Cimitile, and A. De Lucia, “Migrating legacy systems to the web: an experience report,” in *Software Maintenance and Reengineering, 2001. Fifth European Conference on*, pp. 148–157, IEEE, 2001.
- [28] E. Stroulia, M. El-Ramly, P. Iglinski, and P. Sorenson, “User interface reverse engineering in support of interface migration to the web,” *Automated Software Engineering*, vol. 10, no. 3, pp. 271–301, 2003.

Acronyms

DDD domain-driven design. 10, 11

GUI graphical user interface. 5, 6, 9, 12

MVC model-view-controller. 9

REST representational state transfer. 10

Application manual

A.1 Installation

This appendix provides instructions on running the demo application.

A.2 User manual

A concise manual to use the interface.

Time trial test class

Here I will describe how the time trial has been implemented in Java.