

BACHELOR INFORMATICA



UNIVERSITY OF AMSTERDAM

CLIsis: An Interface for Visually Impaired Users of Apache Isis Applications

Sander Ginn

May 30, 2016

Supervisor(s): Maarten Marx (UvA), Dan Haywood (Apache)

Signed: Maarten Marx (UvA)

Abstract

This will be the abstract.

Contents

1	Introduction	4
1.1	Research questions	6
1.2	Overview of thesis	6
2	Theoretical background	7
2.1	Apache Isis	9
2.1.1	Domain-driven development	9
2.1.2	Naked Objects pattern	10
2.1.3	Java Framework	11
2.2	Previous work	11
2.3	Pitfalls in user interface adaption	12
3	Methods	13
3.1	User interface migration	15
3.1.1	Detection	15
3.1.2	Representation	15
3.1.3	Transformation	15
3.2	Experiments	15
3.2.1	GOMS	15
3.2.2	Time trial	15
4	Implementation	16
4.1	Frameworks and tools	17
4.1.1	REST API	17
4.1.2	AngularJS	17
4.1.3	Web Speech API	17
4.2	Functionality	17
4.2.1	Main module	17
4.2.2	Controllers	17
4.2.3	Views	17
5	Evaluation	18
5.1	Specification requirements	19
5.2	Experiments	19
5.2.1	GOMS	19
5.2.2	Time trial	19
5.2.3	Performance difference	19
6	Conclusion	20
6.1	Future work	21
6.2	Acknowledgements	21
A	Application manual	28
A.1	Installation	29
A.2	User manual	29
B	Time trial test class	30

Introduction

Every application that offers a method of interaction with a human being requires some form of user interface to enable this interaction. But what exactly characterises a user interface? Galitz[13] provides the following definition for 'user interface':

The user interface is the part of a computer and its software that people can see, hear, touch, talk to, or otherwise understand or direct.

Thus, in the context of an application, a user interface is a means of interaction between the user and system based on mutually understood communication methods. These communication methods can range from visual methods such as computer displays to audible methods like speakers.

This is still a very abstract definition of what a user interface does. It does not specify how system information is presented, how the user provides input, what permissions and restrictions the interface must adhere to, et cetera. Nielsen[24] describes an overview of user interfaces throughout history which will provide some practical examples.

When the computer was in its infancy, several primitive user interfaces have been developed. As the use of computers became more widespread, however, the *line-oriented interface* was adopted as a method of interaction. The name is derived from the characteristic that the user interacted with a system on a single line; once submitted, the command could not be modified anymore. As long as the information is well structured this method of interaction still has valid use cases in present times, as it effectively delimits functionality, preventing novice users from running into issues. Assistive technology for visually impaired users excelled with line-oriented interfaces; as all content was character based, it could easily be made interpretable for someone who lacks (good) vision, for example through audio[27].

The successor of line-oriented interfaces is the *full-screen interface*. This interface aims to exploit a greater deal of the screen real estate by enabling the user to interact with more elements than just one input line. The full-screen interface also introduced the concept of menus, where functionality available to the user is listed in a hierarchical structure. Menu design has been researched extensively[25, 17, 11] to determine the optimal balance of depth and breadth. Depth decreases complexity of menus while increasing ease of navigation; the contrary applies to breadth. However, menu-based interfaces also marked the beginning of a decrease in accessibility for visually impaired users, as the information displayed gradually became more catered towards a visual mental model. Visually impaired users tend to mentally model information differently, and thus directly translating a visual interface to a non-visual representation can be extremely difficult to comprehend for the target audience[8].

Finally, the user interface type which is currently the most widespread, is the *graphical user interface* (GUI). The vast majority of modern applications offer a graphical user interface as their primary method of interaction. A key characteristic of GUIs is that interaction is offered through direct manipulation. Rather than issuing a command which exactly specifies those parameters of what has to happen, the user gets continuous feedback of their input commands. The obvious advantage of a GUI is that it can take advantage of the human sight to represent data in a more intuitive way. Research has shown that when the interface adheres to some fundamental principles, a user's cognitive performance increases with a GUI. An example of one of these principles is Miller's Law, proposed in his seminal research on information processing[21]. However, Nielsen also addresses the issue that motivates this research: a GUI sidelines the visually impaired from using an application.

Enabling the visually impaired to use applications despite a lack of or very bad vision has been an ongoing effort since GUIs were widely accepted as the de facto standard[7]. However, it has proven to be difficult to design a GUI that is also accessible without visual cues. A big issue is that standardisation has proven ineffective: even though the United States has passed laws that enforce accessibility of websites[4] and the World Wide Web Consortium has developed the Web Content Accessibility Guidelines[6], research has shown that these guidelines are very often ignored and thus websites are rendered inaccessible for accessibility tools such as screenreaders[20].

Visually impaired computer users often make use of a *screen reader* to translate visual objects into a representation that can be interpreted by assistive technology, such as braille terminals or text-to-speech. However, Lazar et al. found that 4 out of 5 top causes of frustration for blind users of screen readers were related to poor design choices of the software or website developer[18]. As a solution to these problems, we will implement a bespoke interface that does not require the use of a third party screen reader.

This research aims to provide an alternative interface for visually impaired users for applications developed with the Apache Isis[1] framework. This interface will not serve as a replacement for the existing interface nor blend in with it. The interface is derived from the metamodel of the framework and thus content independent. This means that the interface is readily available for any Isis application and no alterations are necessary to add it to an existing application.

1.1 Research questions

The central research question in this thesis is:

Can a content independent graphical user interface be adapted to a non-visual interface so that visually impaired users are able to interact with the application?

For the non-visual interface to be useful it is imperative that it offers the same functionality as the graphical user interface. Therefore, we will also answer the following subquestion:

When implementing a non-visual interface, can the integrity of the domain model be maintained while providing an effective and simple method of interaction?

Finally, it is desired that the non-visual interface does not severely impact user performance, as this would imply that it is not a legitimate alternative to the existing interface. Thus, a second subquestion will be answered:

Compared to the standard interface, how is performance affected when the user employs the new interface?

1.2 Overview of thesis

A short overview of what this thesis will describe is given.

Theoretical background

User interface design is a thoroughly studied discipline with strong roots in psychology. In the 1980s GUI development exploded due to better hardware[22]. This meant that traditional user interfaces had to be redesigned to accommodate to the graphical features of the modern computer. Chapter 1 described the issues visually impaired users experienced with this new method of interaction. Section 2.2 expands on some existing efforts on improving usability for the visually impaired. Furthermore, section 2.1 will describe what Apache Isis entails.

2.1 Apache Isis

As a software engineer, picking a framework for developing web applications can be a tedious process. There are dozens of frameworks for Java alone, with the oldest and most adopted one being Spring[5]. The vast majority of these frameworks are based on the *model-view-controller* (MVC) pattern, where the view displays information to the user, the controller processes interaction between the user and the view, and the model contains the information and logic that manipulates this information[19]. The relations between the components are depicted in figure 2.1.




Figure 2.1: Interaction between the components of an MVC application

While the MVC pattern itself has a lot of advantages, it has received criticism in the context of web development. The concept of separating business logic from presentation logic is often not adhered to in web applications, resulting in controllers that are contaminated with logic that should live in the model[3]. Trygve Reenskaug, who introduced the MVC pattern while working at the Xerox Palo Alto Research Center in 1978, concluded that true MVC-based applications were hard to achieve[28]:

”The conventional wisdom in the group was that objects should be visible and tangible, thus bridging the gap between the human brain and the abstract data within the computer. This simple and powerful idea failed ... because users were used to seeing objects from different perspectives. The visible and tangible object would get very complex if it should be able to show itself and be manipulated in many different ways.”

2.1.1 Domain-driven development

Even though Reenskaug did not consider the MVC pattern to be successful, the philosophy of displaying information in a manner that feels natural to humans is still a very valid one. Many

methodologies have been developed in an attempt to capture this philosophy, of which *domain-driven design* (DDD) is a well-known example. Conceived by Eric Evans in 2004, DDD strives to streamline and align the domain experts (the business experts) and the technical experts (the developers). By clearly defining what terms to use for what concepts and ensuring that both parties understand what they mean, a *ubiquitous language* is formed, and information asymmetry can be avoided[10]. The *domain model* is then built solely with terms in the ubiquitous language. The key to success is then to apply model-driven development; by developing the software while strictly adhering to the specified domain model, all stakeholders are able to comprehend the content of the software. When the domain model proves to be insufficient to implement a new feature, the stakeholders should work together to extend the ubiquitous language so that the domain model becomes complete again. This creates a bidirectional relationship between the domain model and the code: when the domain model changes, so does the code and vice versa.

2.1.2 Naked Objects pattern

One of the frustrations often expressed regarding DDD is that while the ubiquitous language combined with the derived domain model may very well tackle the problem of goal diversion, it also increases overhead greatly. Consider figure 2.2: the domain model is represented in the business logic layer. As the domain model will be refined over the course of time, and that any changes to the domain model must be reflected in the code. This means that any modifications to the domain model will have to be applied to the other layers, too.




Figure 2.2: The main application layers in DDD

To address this issue, Richard Pawson designed the *Naked Objects* (NO) pattern, based on three core concepts[26]:

1. Business logic should only reside in domain objects¹
2. The user interface is an exact representation of the domain objects, and thus is an object-oriented user interface
3. The user interface is derived completely and fully automatically from the domain objects through reflection²

These three concepts combined provide a means of developing software that robustly follows the philosophy of DDD. The first two concepts ensure that the software implements the domain model and prevents the use of technical language. The application, whether prototypal or deployed, will feel familiar to the stakeholders and developers alike. The third concept is not necessarily related to DDD, but does take away the burden of developing the presentation layer and thus places the focus on the business logic: i.e., the domain model.

In the foreword of Pawson's thesis, Reenskaug states that the NO pattern could be seen as an important extension and implementation of the original MVC pattern.

¹A domain object is an entity that holds relevant data, in order to transfer the data between the application layers.

²Reflection is a language's ability to inspect and dynamically call classes, methods, attributes, etc. at runtime.

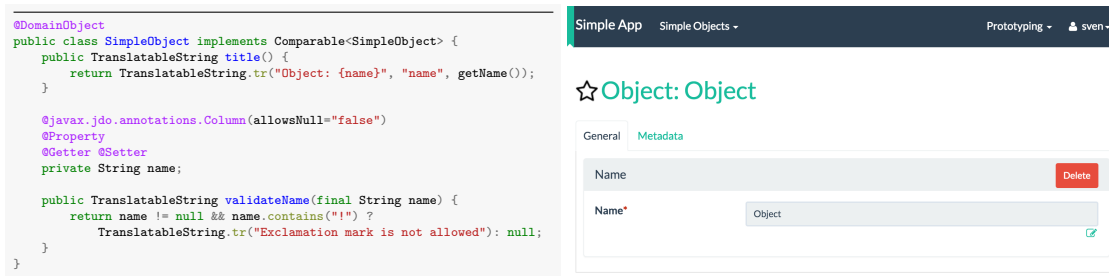


Figure 2.3: A very simple example of how the user interface is derived from code

2.1.3 Java Framework

Apache Isis is a framework written in Java that implements the NO pattern. The core strength of the framework is that behaviour which would normally be defined in controllers in classical MVC applications is automatically generated by reflection. A major benefit of this feature is that prototyping is greatly simplified; the business logic is virtually all that needs to be programmed, and the framework does the rest. This means that project managers, product owners and other parties involved can be presented with working prototypes rather than charts and specifications, and it relieves developers from wasting precious time on coding the visible part of an application - code which often has limited reusability in a later stage of development.

This does not, however, imply that Apache Isis is merely useful for prototyping. The automatically generated interface is perfectly suitable for a lot of business oriented applications. Estatio[2], an estate management application built on Apache Isis commissioned by a listed real estate investment firm, makes use of the standard wicket viewer. In case a given application demands a bespoke user interface, it can be designed to use the REST API which ships with Apache Isis, which we will use to implement our interface.

Due to the underlying NO paradigm that the user interface is derived from the domain objects, Apache Isis is exceptionally suitable for this research project. Every bit of information that dictates what is visible to the user is found in the code that describes the functionality of the software, and thus we are able to develop a new user interface that retrieves everything it needs from what the developers have written regardless of the interface they intend to use.

2.2 Previous work

User interfaces make use of at least one of the traditional human senses to convey information. It is evident that GUIs are mainly based on exploiting human vision. As visually impaired people lack this sense, they have to interact with software through what remains: the auditory, somatic, gustatory and olfactory senses. The latter two senses are relatively unexplored in terms of human-computer interaction. While there have been developments such as 'smicons', representing icons through odour[15], and a display that represents rudimentary information through scents sprayed in the direction of the nose[30], it has proven difficult to employ these senses in a practical manner. The main reason is that smell and taste are chemical rather than physical sensations and thus a lot harder to synthesise, let alone accurately[16]. Therefore, effectively two senses remain available: the auditory and somatic senses.

Exploiting the hearing sense is a good way of making software usable for a large audience, as it only requires a form of audio output to function and virtually all modern devices have either built-in speakers or a headphone jack. Most auditory user interfaces synthesise text to speech and use auditory icons for basic operations, such as the sound of crumpling paper on a delete operation[23]. While this is regarded as one of the easiest ways of increasing accessibility of software, it still leaves a portion of the auditory sense unused, as it does not take advantage of spatial sound. Research has shown that basic elements of interfaces, such as menus, can be displayed very effectively in a spatial setting[12].

Software can also interact with the somatic sense through haptic feedback. Tactile feedback is a form of haptic feedback that relies on the sensation of touch, vibration and pressure. Visually impaired people who have learned braille can take advantage of output devices called braille terminals combined with screen reader software, which translates text on the screen to braille. This does not, however, offer practical support for graphical elements. With the advent of mobile technology, research in the field of tactile user interfaces has increased. The majority of these mobile devices have touch screens and vibrating motors, allowing gesture-based input and haptic feedback[14, 31].

2.3 Pitfalls in user interface adaption

The concept of universally accessible user interfaces has gained a lot of traction since the widespread adoption of GUIs. It has proven difficult, however, to adapt an interface in such a way that it becomes both universally accessible and maintainable. The foremost shortcoming of adapting an existing user interface is that it is mainly a posteriori, and thus any future changes to the software will require modifications to all existing user interfaces. Another major issue is that it might be very difficult or even impossible to adapt certain functionality due to the constraints of the new interface[9]. In the present time, it is generally accepted that the proactive approach of making design choices with universal accessibility in mind is more convenient and successful than the reactive approach of adapting an existing design[29]. As this is not within the scope of this research, we will not address this in detail.

Methods

We have selected several methods for implementing the new interface. First, section 3.1 will describe the method applied to the existing interface in order to form a theoretical basis for the implementation. Section 3.2 then outlines what experiment methods are used to evaluate the interface.

3.1 User interface migration

Section ?? expanded on user interface migration in history. With new technologies constantly emerging, user interface design is a continuous process that will always be relevant and extensively studied.

3.1.1 Detection

Step one of the migration method.

3.1.2 Representation

Step two of the migration method.

Specifications

The set of specifications that are derived from section 3.1.1 are presented.

3.1.3 Transformation

Step three of the migration method.

3.2 Experiments

A description of the methods used to conduct the experiments.

3.2.1 GOMS

The Goals, Operators, Methods, Selection rules method is explained.

3.2.2 Time trial

A custom time trial method is described.

Implementation

This chapter describes the practical side of the project.

4.1 Frameworks and tools

This section expands on what available software tools I have used to implement the new interface.

4.1.1 REST API

This subsection will describe how Apache Isis' REST API works and how it is incorporated in the new interface.

4.1.2 AngularJS

Here I will explain the choice for AngularJS.

4.1.3 Web Speech API

This subsection will describe how the Web Speech API is used.

4.2 Functionality

This section will expand on the individual components that process the information from the REST API.

4.2.1 Main module

This subsection describes the main app module.

4.2.2 Controllers

This subsection describes each controller's functionality.

4.2.3 Views

This subsection describes the purpose of each view.

Evaluation

This chapter describes the experiment results and answers the research questions.

5.1 Specification requirements

In this section, I will motivate which specifications have been met in the implementation. With this information I will be able to answer the first subquestion.

5.2 Experiments

5.2.1 GOMS

Here, the results from applying GOMS will be evaluated.

5.2.2 Time trial

Here, the results from the time trial experiment will be evaluated.

5.2.3 Performance difference

This subsection will evaluate the experiment results and answer the second subquestion.

Conclusion

Here, I will answer the research question and make any other conclusions.

6.1 Future work

Any remaining future work can be discussed here.

6.2 Acknowledgements

Bibliography

- [1] Apache isis. <http://isis.apache.org/>. (Accessed on 05/16/2016).
- [2] Estatio. <http://www.estatio.org/>. (Accessed on 05/18/2016).
- [3] Fulfilling the promise of mvc. <https://www.infoq.com/articles/Nacked-MVC>. (Accessed on 05/17/2016).
- [4] Section 508 of the rehabilitation act. <http://www.section508.gov/section-508-of-the-rehabilitation-act>. (Accessed on 05/16/2016).
- [5] Spring. <https://spring.io/>. (Accessed on 05/16/2016).
- [6] Web content accessibility guidelines (wcag) 2.0. <https://www.w3.org/TR/WCAG20/>. (Accessed on 05/16/2016).
- [7] Boyd, L. H. et al. (1990). The graphical user interface crisis: Danger and opportunity.
- [8] Edwards, W. K., Mynatt, E. D., and Stockton, K. (1994). Providing access to graphical user interfaces not graphical screens. In *Proceedings of the first annual ACM conference on Assistive technologies*, pages 47–54. ACM.
- [9] Emiliani, P. L. and Stephanidis, C. (2000). From adaptations to user interfaces for all. In *6th ERCIM workshop CNR-IROE, Florence, Italy*.
- [10] Evans, E. (2004). *Domain-driven design: tackling complexity in the heart of software*. Addison-Wesley Professional.
- [11] Fisher, D. L., Yungkurth, E. J., and Moss, S. M. (1990). Optimal menu hierarchy design: syntax and semantics. *Human Factors: The Journal of the Human Factors and Ergonomics Society*, 32(6):665–683.
- [12] Frauenberger, C., Putz, V., and Holdrich, R. (2004). Spatial auditory displays—a study on the use of virtual audio environments as interfaces for users with visual disabilities. *DAFx04 Proceedings*, pages 5–8.
- [13] Galitz, W. O. (2007). *The essential guide to user interface design: an introduction to GUI design principles and techniques*. John Wiley & Sons.
- [14] Kane, S. K., Bigham, J. P., and Wobbrock, J. O. (2008). Slide rule: making mobile touch screens accessible to blind people using multi-touch interaction techniques. In *Proceedings of the 10th international ACM SIGACCESS conference on Computers and accessibility*, pages 73–80. ACM.
- [15] Kaye, J. N. (2001). *Symbolic olfactory display*. PhD thesis, Citeseer.
- [16] Kortum, P. (2008). *HCI beyond the GUI: Design for haptic, speech, olfactory, and other nontraditional interfaces*. Morgan Kaufmann.
- [17] Landauer, T. K. and Nachbar, D. (1985). Selection from alphabetic and numeric menu trees using a touch screen: breadth, depth, and width. *ACM SIGCHI Bulletin*, 16(4):73–78.
- [18] Lazar, J., Allen, A., Kleinman, J., and Malarkey, C. (2007). What frustrates screen reader users on the web: A study of 100 blind users. *International Journal of human-computer interaction*, 22(3):247–269.

- [19] Leff, A. and Rayfield, J. T. (2001). Web-application development using the model/view/controller design pattern. In *Enterprise Distributed Object Computing Conference, 2001. EDOC'01. Proceedings. Fifth IEEE International*, pages 118–127. IEEE.
- [20] Leuthold, S., Bargas-Avila, J. A., and Opwis, K. (2008). Beyond web content accessibility guidelines: Design of enhanced text user interfaces for blind internet users. *International Journal of Human-Computer Studies*, 66(4):257–270.
- [21] Miller, G. A. (1956). The magical number seven, plus or minus two: some limits on our capacity for processing information. *Psychological review*, 63(2):81.
- [22] Myers, B. A. (1998). A brief history of human-computer interaction technology. *interactions*, 5(2):44–54.
- [23] Mynatt, E. D. (1995). Transforming graphical interfaces into auditory interfaces. In *Conference companion on Human factors in computing systems*, pages 67–68. ACM.
- [24] Nielsen, J. (1994). *Usability engineering*. Elsevier.
- [25] Paap, K. R. and Roske-Hofstrand, R. J. (1986). The optimal number of menu options per panel. *Human Factors: The Journal of the Human Factors and Ergonomics Society*, 28(4):377–385.
- [26] Pawson, R. and Matthews, R. (2002). Naked objects. In *Companion of the 17th annual ACM SIGPLAN conference on Object-oriented programming, systems, languages, and applications*, pages 36–37. ACM.
- [27] Poll, L. H. D. (1996). *Visualising Graphical User Interfaces for Blindusers*. Eindhoven University.
- [28] Reenskaug, T. (2007). Programming with roles and classes: The babyuml approach. *A chapter in Computer Software Engineering Research*.
- [29] Savidis, A. and Stephanidis, C. (2004). Unified user interface design: designing universally accessible interactions. *Interacting with computers*, 16(2):243–270.
- [30] Yanagida, Y., Kawato, S., Noma, H., Tomono, A., and Tesutani, N. (2004). Projection based olfactory display with nose tracking. In *Virtual Reality, 2004. Proceedings. IEEE*, pages 43–50. IEEE.
- [31] Yatani, K. and Truong, K. N. (2009). Semfeel: a user interface with semantic tactile feedback for mobile touch-screen devices. In *Proceedings of the 22nd annual ACM symposium on User interface software and technology*, pages 111–120. ACM.

Acronyms

DDD domain-driven design. 10

GUI graphical user interface. 5, 6, 9, 11

MVC model-view-controller. 9–11

NO Naked Objects. 10, 11

REST representational state transfer. 11

Application manual

A.1 Installation

This appendix provides instructions on running the demo application.

A.2 User manual

A concise manual to use the interface.

Time trial test class

Here I will describe how the time trial has been implemented in Java.